


Neuroengineering (I)

6. Encoder-Decoder networks

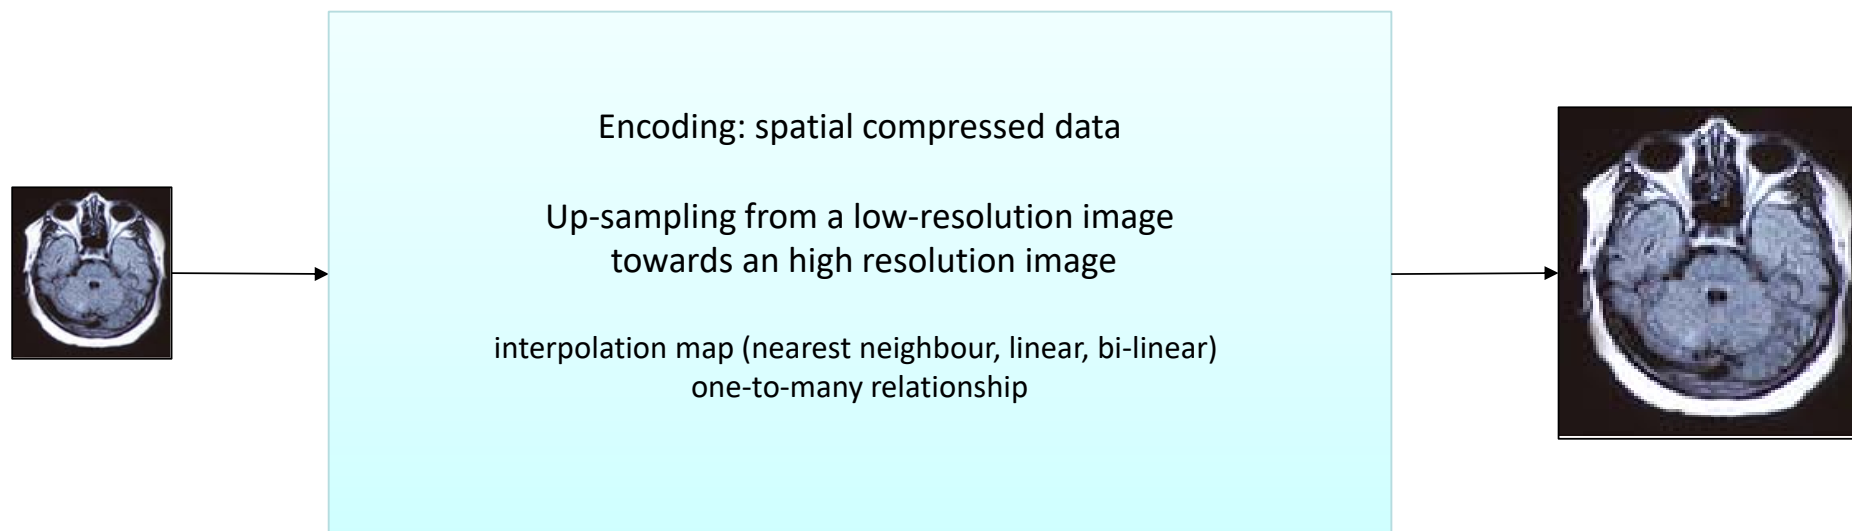
- **Scuola di Ingegneria Industriale e dell'Informazione**
– Politecnico di Milano
- Prof. Pietro Cerveri

Generative process

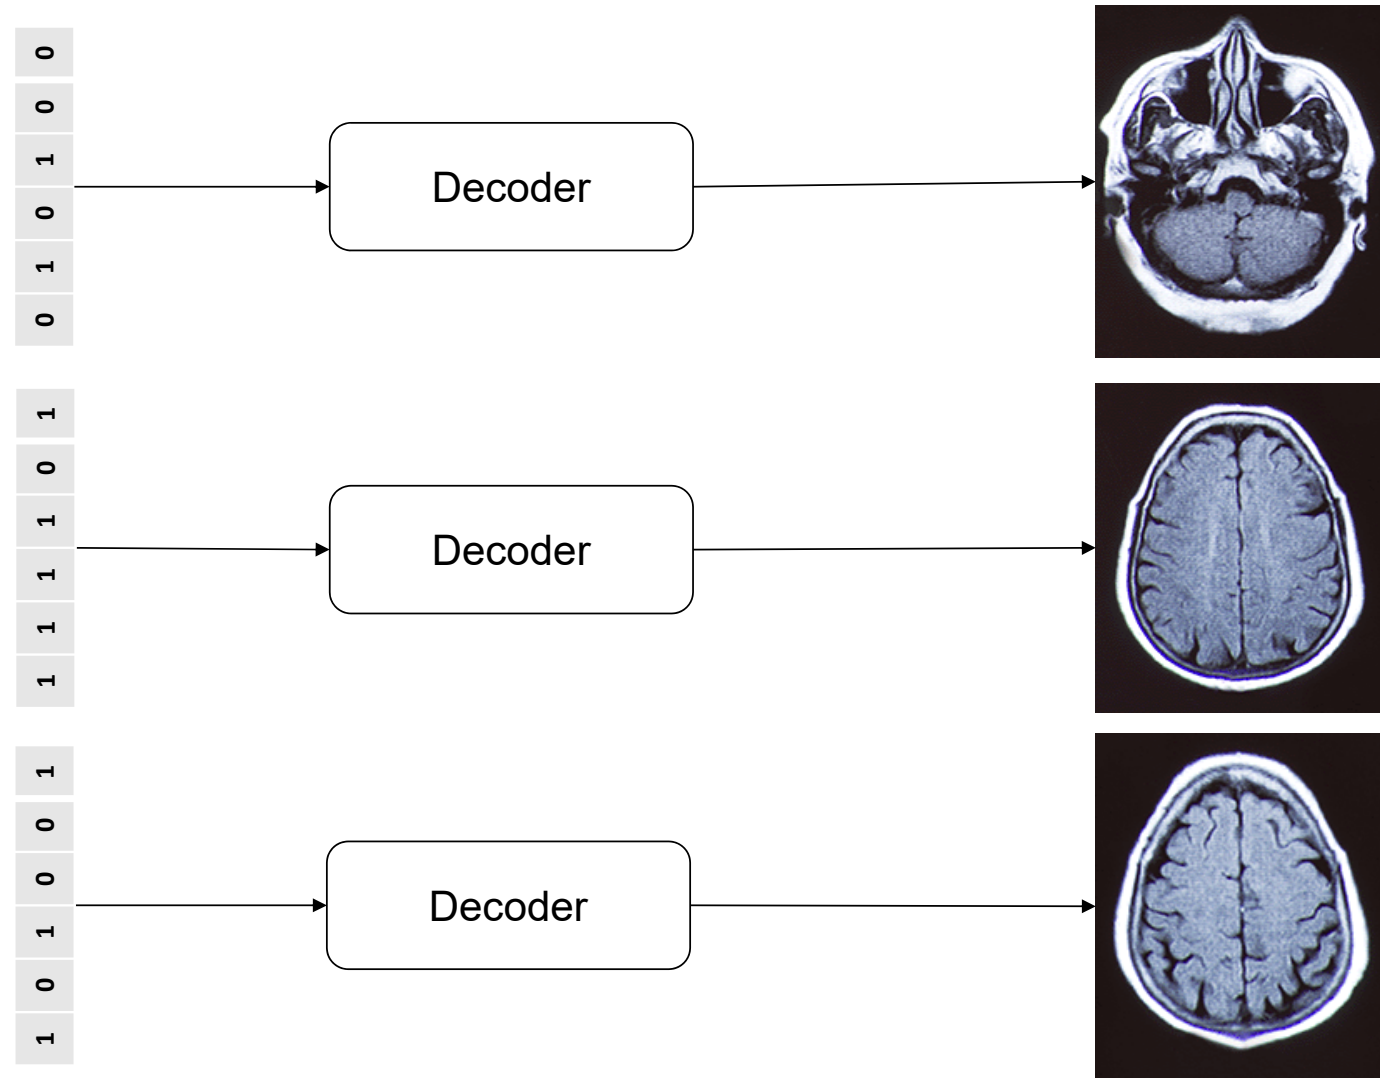


- Process of decoding the information represented into a vector whose elements are extracted from predefined “alphabet”
- How uncompressing/decode the encoded information?
 - Complex information requires appropriate transforms
 - Single (interpolation) against progressive decoding (multiple sequential transforms)


Up-sampling as a decoding process



Encoding alphabet



AN networks for decoding data



- Deep learning approach to optimal up-sampling: **transposed convolution** (deconvolution) which does not use a predefined interpolation method
- Learning will allow to tuning of the parameters (weights of the neurons)

Filter			
-1	1	-1	
1	2	1	
-1	1	-1	
Input image			
2	2	1	1
2	2	3	1
4	5	1	1
4	1	0	4

Revise 2D convolution

3x3 filter

4x4 input image

Stride length: 1

No padding

Output size??

2D Convolution, no padding, stride: 1

The diagram illustrates the process of applying a 3x3 filter to a 3x3 input image to produce a 2x2 output. The filter values are consistently [-1, 1, -1] across all stages. The input images and their corresponding outputs are as follows:

- Stage 1:** Input image $\begin{bmatrix} 2 & 2 & 1 \\ 2 & 2 & 3 \\ 4 & 5 & 1 \end{bmatrix}$, Output $\begin{bmatrix} 8 & 0 \\ 0 & 0 \end{bmatrix}$.
- Stage 2:** Input image $\begin{bmatrix} 2 & 2 & 1 \\ 2 & 2 & 3 \\ 4 & 5 & 1 \end{bmatrix}$, Output $\begin{bmatrix} 8 & 2 \\ 0 & 0 \end{bmatrix}$.
- Stage 3:** Input image $\begin{bmatrix} 2 & 2 & 1 \\ 2 & 2 & 3 \\ 4 & 5 & 1 \end{bmatrix}$, Output $\begin{bmatrix} 8 & 2 \\ 9 & 0 \end{bmatrix}$.
- Stage 4:** Input image $\begin{bmatrix} 2 & 2 & 1 \\ 2 & 2 & 3 \\ 4 & 5 & 1 \end{bmatrix}$, Output $\begin{bmatrix} 8 & 2 \\ 9 & 4 \end{bmatrix}$.

the 3×3 filter is used to connect the 9 values in the 4×4 input matrix to 1 value in the 2×2 output matrix

2D Convolution, no padding, stride: 1

The diagram illustrates four stages of a 2D convolution operation. Each stage shows a 3x3 input image, a 3x3 filter, and a 2x2 output. The filter is consistently $\begin{bmatrix} -1 & 1 & -1 \\ 1 & 2 & 1 \\ -1 & 1 & -1 \end{bmatrix}$. The input image is consistently $\begin{bmatrix} 2 & 2 & 1 & 1 \\ 2 & 2 & 3 & 1 \\ 4 & 5 & 1 & 1 \\ 4 & 1 & 0 & 4 \end{bmatrix}$. The output is consistently $\begin{bmatrix} 8 & 2 \\ 9 & 4 \end{bmatrix}$.

the 3×3 filter is used to connect the 9 values in the 4×4 input matrix to 1 value in the 2×2 output matrix

Going backward (pixel-wise multiplication, stride 1)

The diagram illustrates three stages of a 2D convolution operation. Each stage shows a 3x3 input, a 3x3 filter, and a 4x4 output image.

Stage 1:

- Filter:**

$$\begin{bmatrix} -1 & 1 & -1 \\ 1 & 2 & 1 \\ -1 & 1 & -1 \end{bmatrix}$$
- Input:**

$$\begin{bmatrix} 8 & 2 \\ 9 & 4 \end{bmatrix}$$
- Output image:**

$$\begin{bmatrix} -8 & 8 & -8 & 0 \\ 8 & 16 & 8 & 0 \\ -8 & 8 & -8 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

Stage 2:

- Filter:**

$$\begin{bmatrix} -1 & 1 & -1 \\ 1 & 2 & 1 \\ -1 & 1 & -1 \end{bmatrix}$$
- Input:**

$$\begin{bmatrix} 8 & 2 \\ 9 & 4 \end{bmatrix}$$
- Output image:**

$$\begin{bmatrix} -2 & 2 & -2 & 0 \\ 2 & 4 & 2 & 0 \\ -2 & 2 & -2 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

Stage 3:

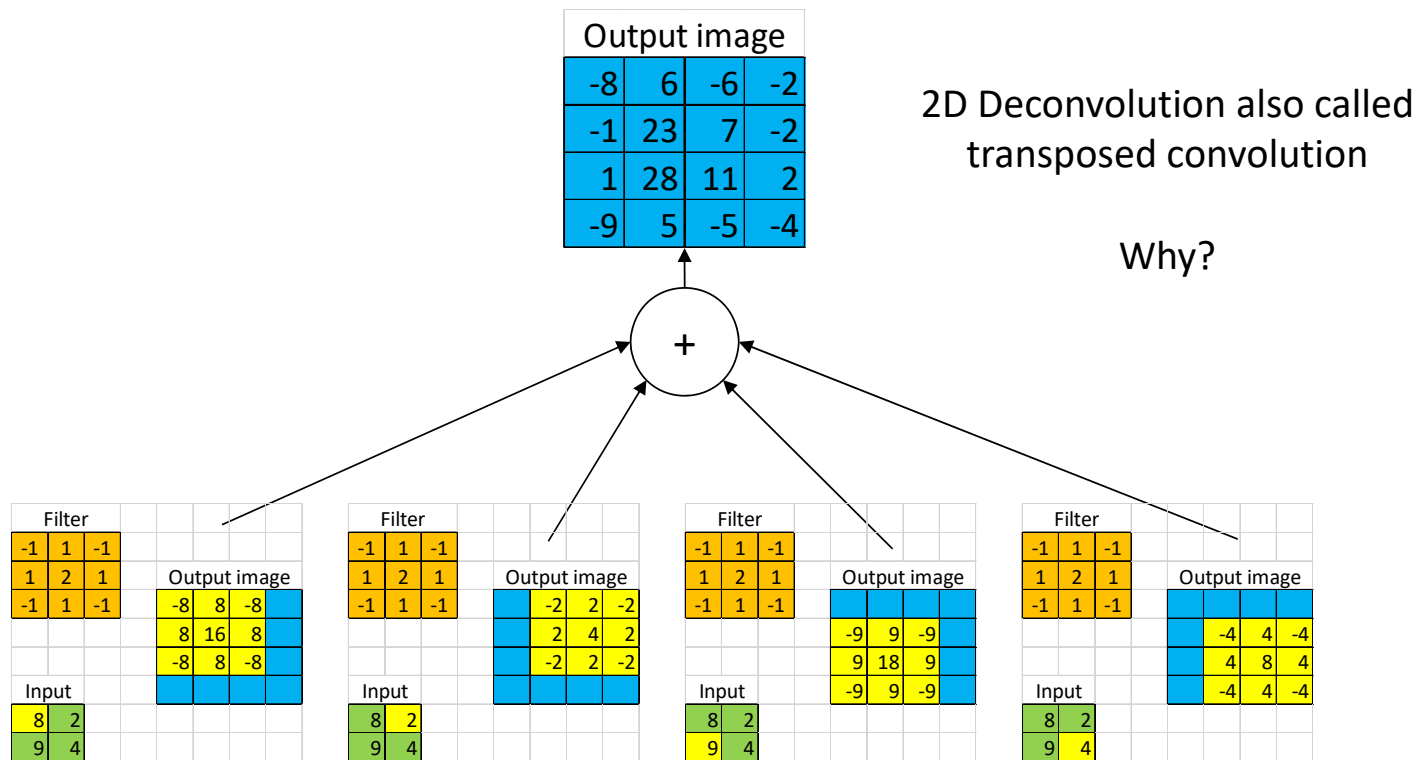
- Filter:**

$$\begin{bmatrix} -1 & 1 & -1 \\ 1 & 2 & 1 \\ -1 & 1 & -1 \end{bmatrix}$$
- Input:**

$$\begin{bmatrix} 8 & 2 \\ 9 & 4 \end{bmatrix}$$
- Output image:**


$$\begin{bmatrix} -9 & 9 & -9 & 0 \\ 9 & 18 & 9 & 0 \\ -9 & 9 & -9 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

the 3x3 filter is used to up-sampling the 2x2 input matrix by pixel-wise multiplication, kernel striding and summation




the 3x3 filter is used to up-sampling the 2x2 input matrix by pixel-wise multiplication, kernel striding and summation

Matrix convolution



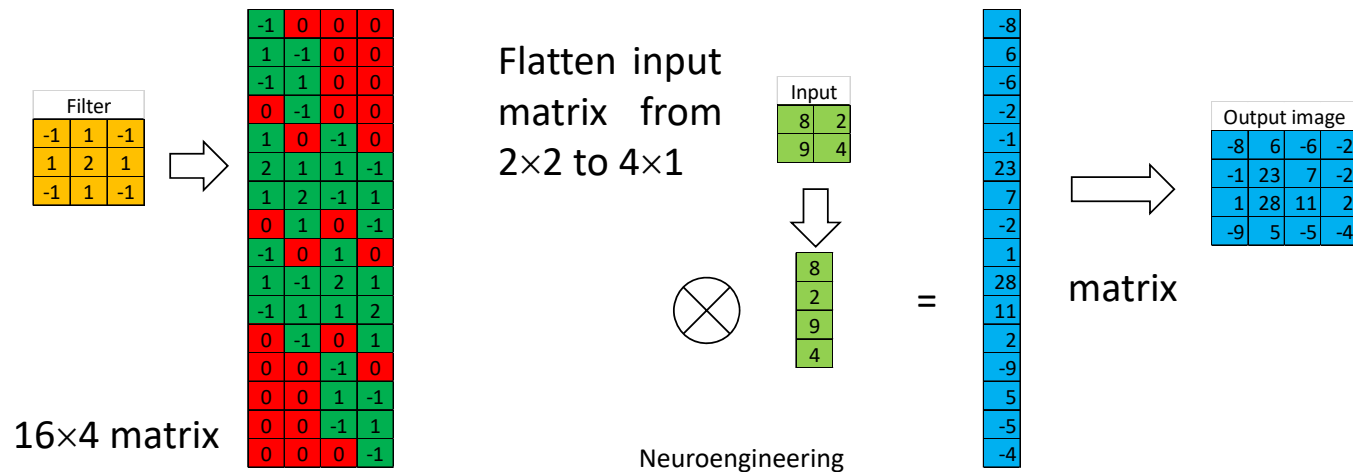
Rearranged filter weights according to size and the stride

Matrix convolution



Rearranged filter weights according to size and the stride

Matrix deconvolution (transposed convolution)




Remarks

- The transposed convolution operation forms the same connectivity as the normal convolution but in the backward direction
 - ❑ one-to-many rather than many-to-one association
- As such, the transposed convolution is not a convolution. But we can emulate the transposed convolution using a convolution. We up-sample the input by adding zeros between the values in the input matrix in a way that the direct convolution produces the same effect as the transposed convolution.
- We can use it to conduct up-sampling. Moreover, the **weights in the transposed convolution are learnable**. So **we do not need a predefined interpolation method**.


Deconvolution by zero-padding the input image

A deconvolution for one 3×3 filter with stride 1 and image (blue entries) with zero-padding of 2 (white entries)



$$s_x = \frac{(I_x + 2z_x - f_x)}{L_x} + 1$$
$$s_y = \frac{(I_y + 2z_y - f_y)}{L_y} + 1$$

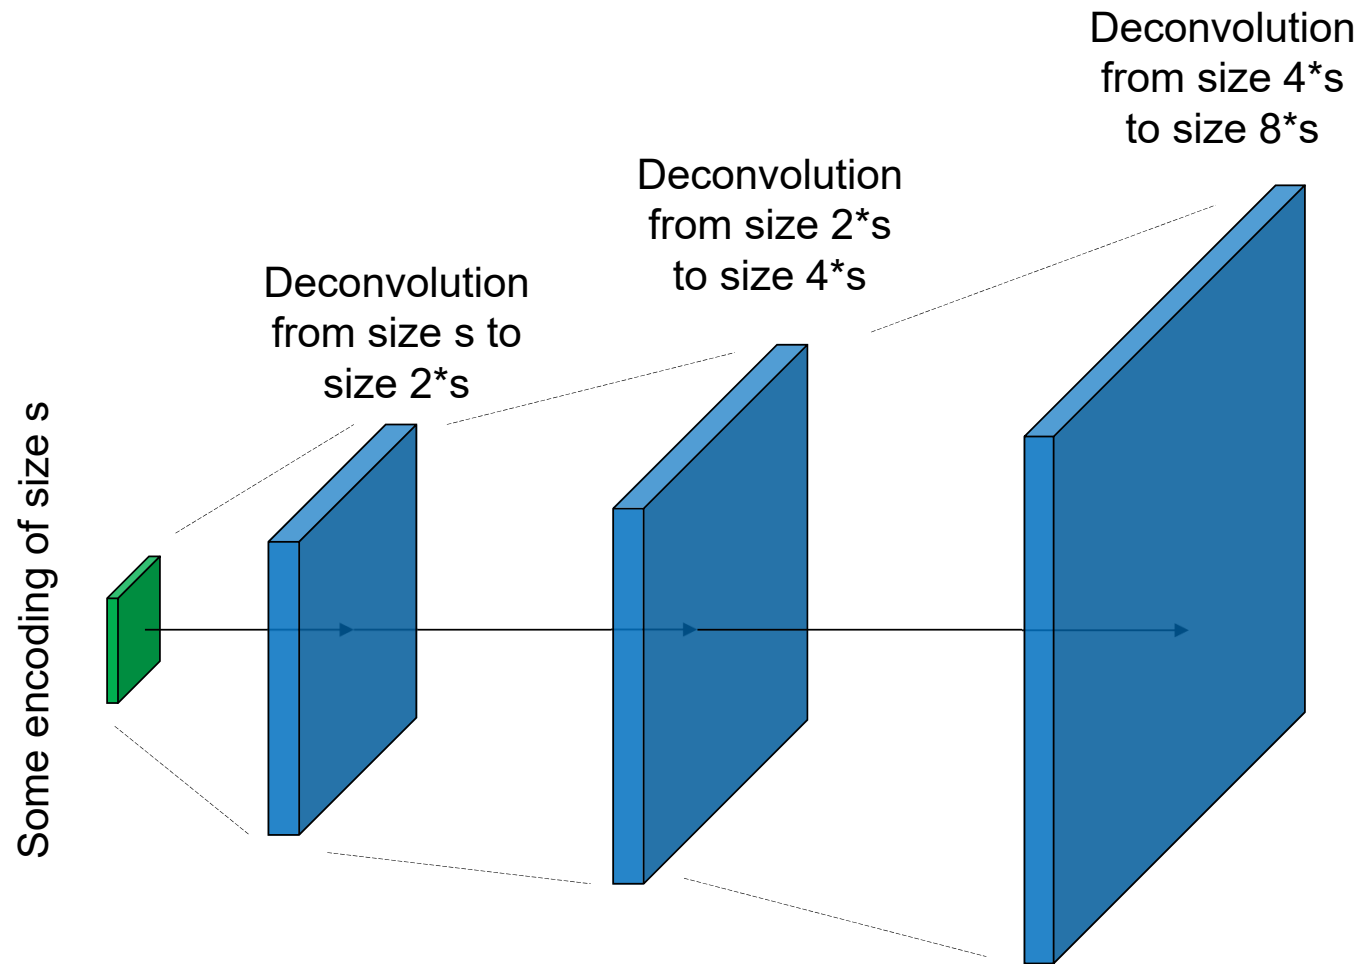
Inter-pixel zero-padding of 1, the deconvolution would look like this
Also called **fractional strided convolution**



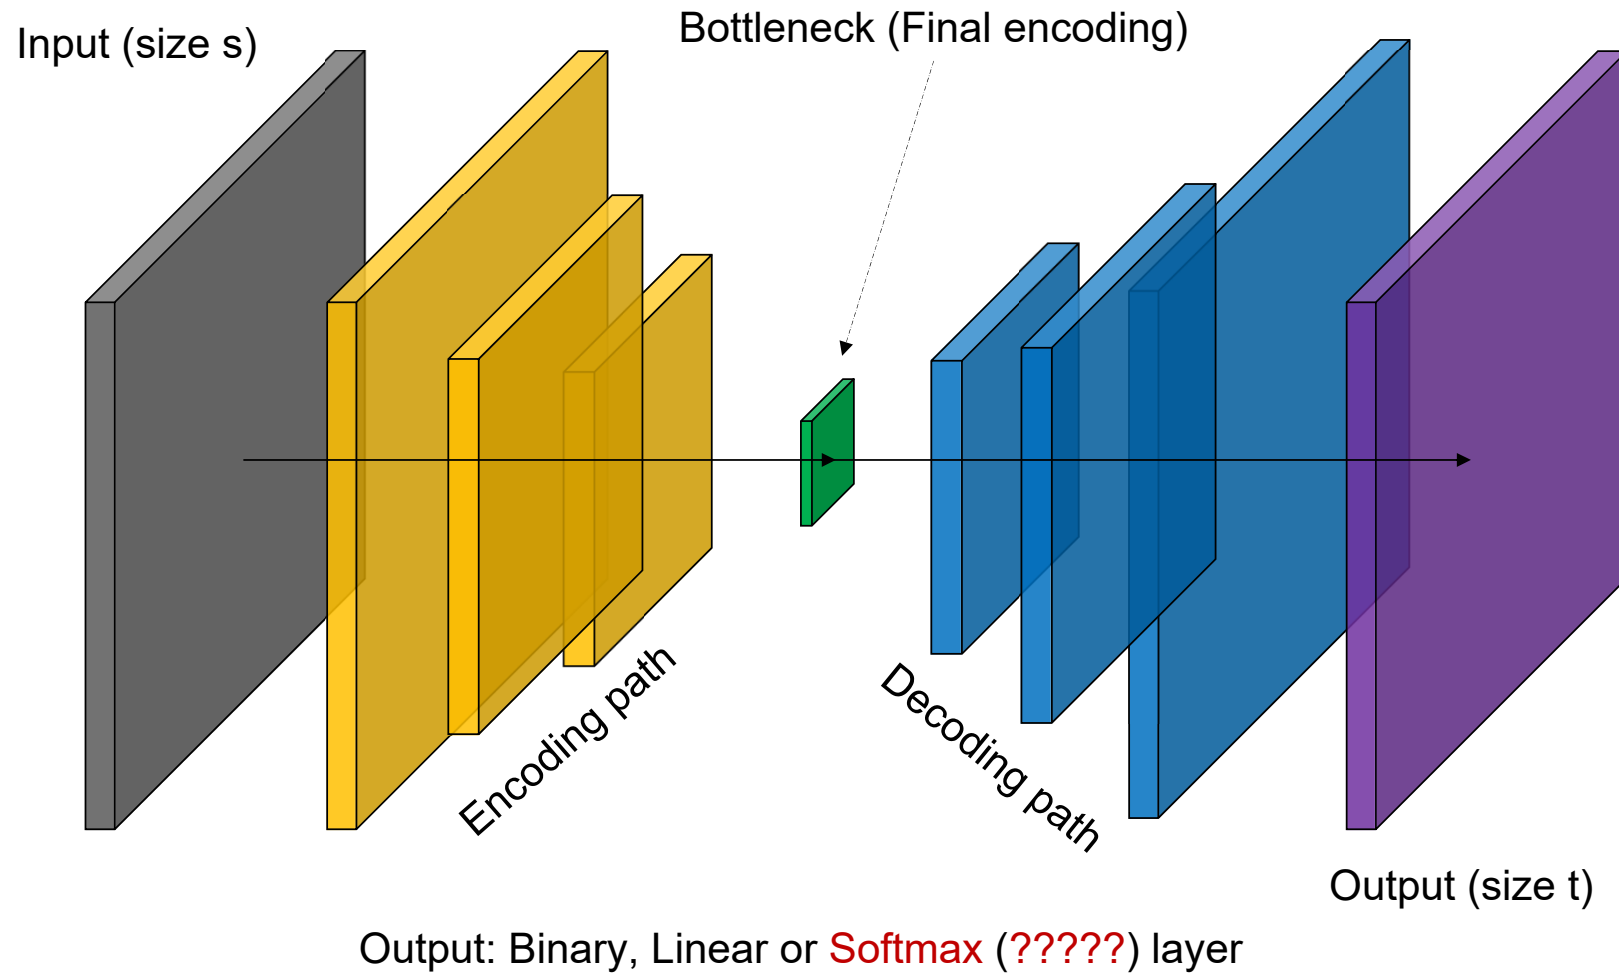
$$s_x = \frac{(I_x + 4z_x - f_x)}{L_x} + 1$$
$$s_y = \frac{(I_y + 4z_y - f_y)}{L_y} + 1$$

Upsampling convolution

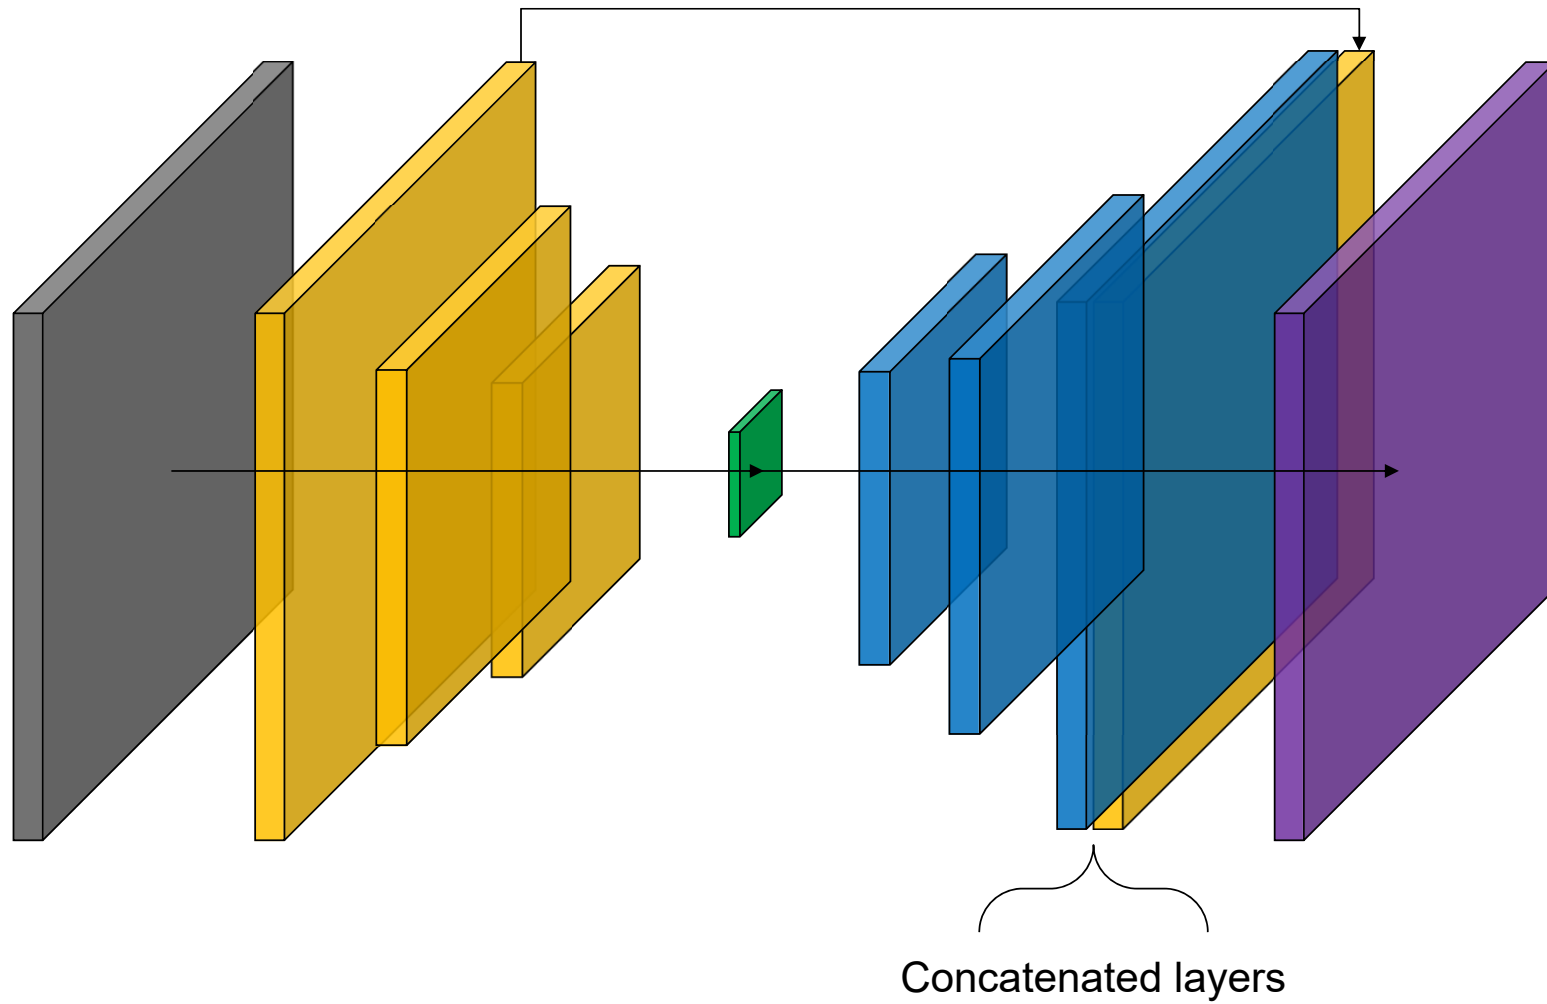
Deconvolution networks




Encoder-decoder network




Unet -> Encoder-decoder network with skip connections



Concatenation between two or more layers




Graphical symbol for concatenation



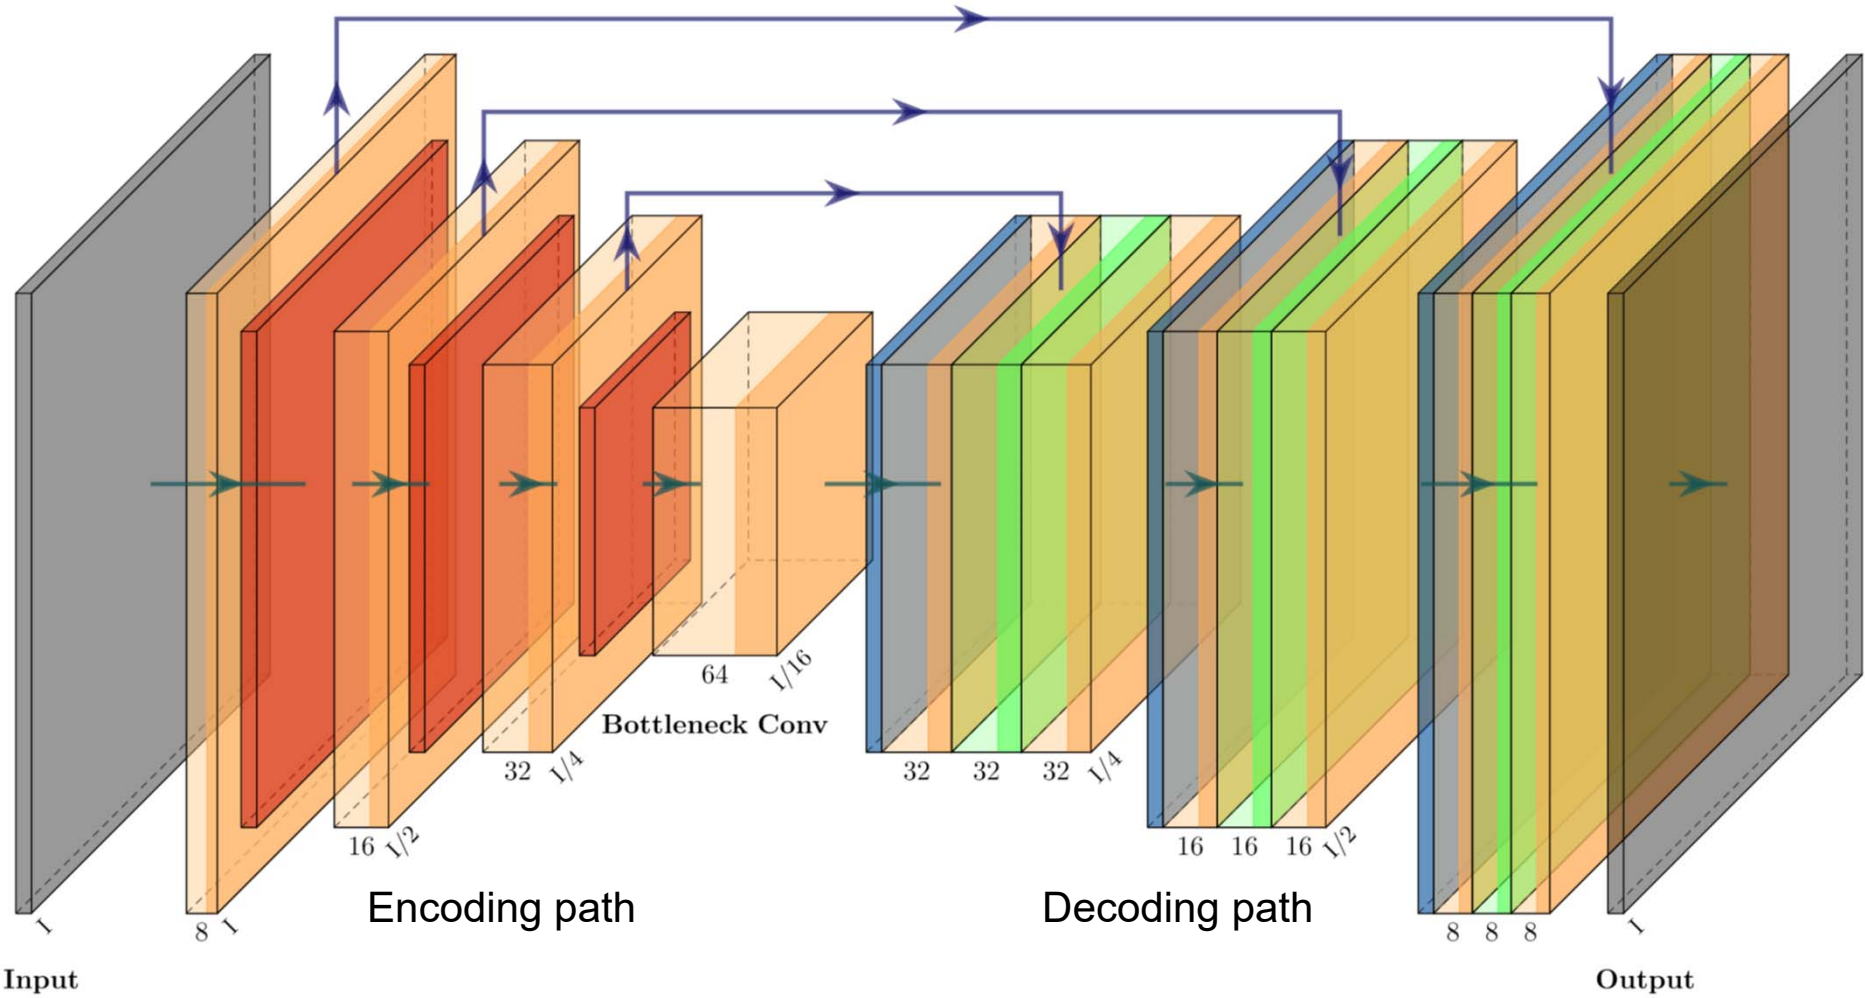
All concatenated layers must have the same size

Role of the **skip connection**



Using the encoded information at the same processing level to aid the up-sampling

General Unet architecture



Encoding layer

Convolution (i is the layer index)

```
X = Conv3D(nFM_i, kernel_size = (3, 3, 3), strides = (1, 1, 1), padding = 'same', activation = 'linear', )(X_from previous layer)
```

ReLu

```
X = Activation('relu')(X)
```

Skip connection

```
Li_X = X
```

Maxpooling

```
X_tonextlayer = MaxPooling3D(pool_size = (2, 2, 2), strides = (2, 2, 2), padding = 'same')(X)
```

Decoding layer

Deconvolution (j is the layer index)

```
X = Conv3DTranspose(nFM_j, kernel_size = (2, 2, 2), strides = (2, 2, 2), padding = 'same', activation = 'linear', )(X_from  
previous layer)
```

ReLu

```
X = Activation('relu')(X)
```

Concatenation with Skip connection

```
X = Concatenate(X, Li_X)
```


Convolution (j the layer index)

```
X = Conv3D(nFM_j, kernel_size = (3, 3, 3), strides = (1, 1, 1), padding = 'same', activation = 'linear', )(X)
```

ReLu


```
X_tonext layer = Activation('relu')(X)
```

Another way to look at the Unet



Unet in action: image segmentation


Input

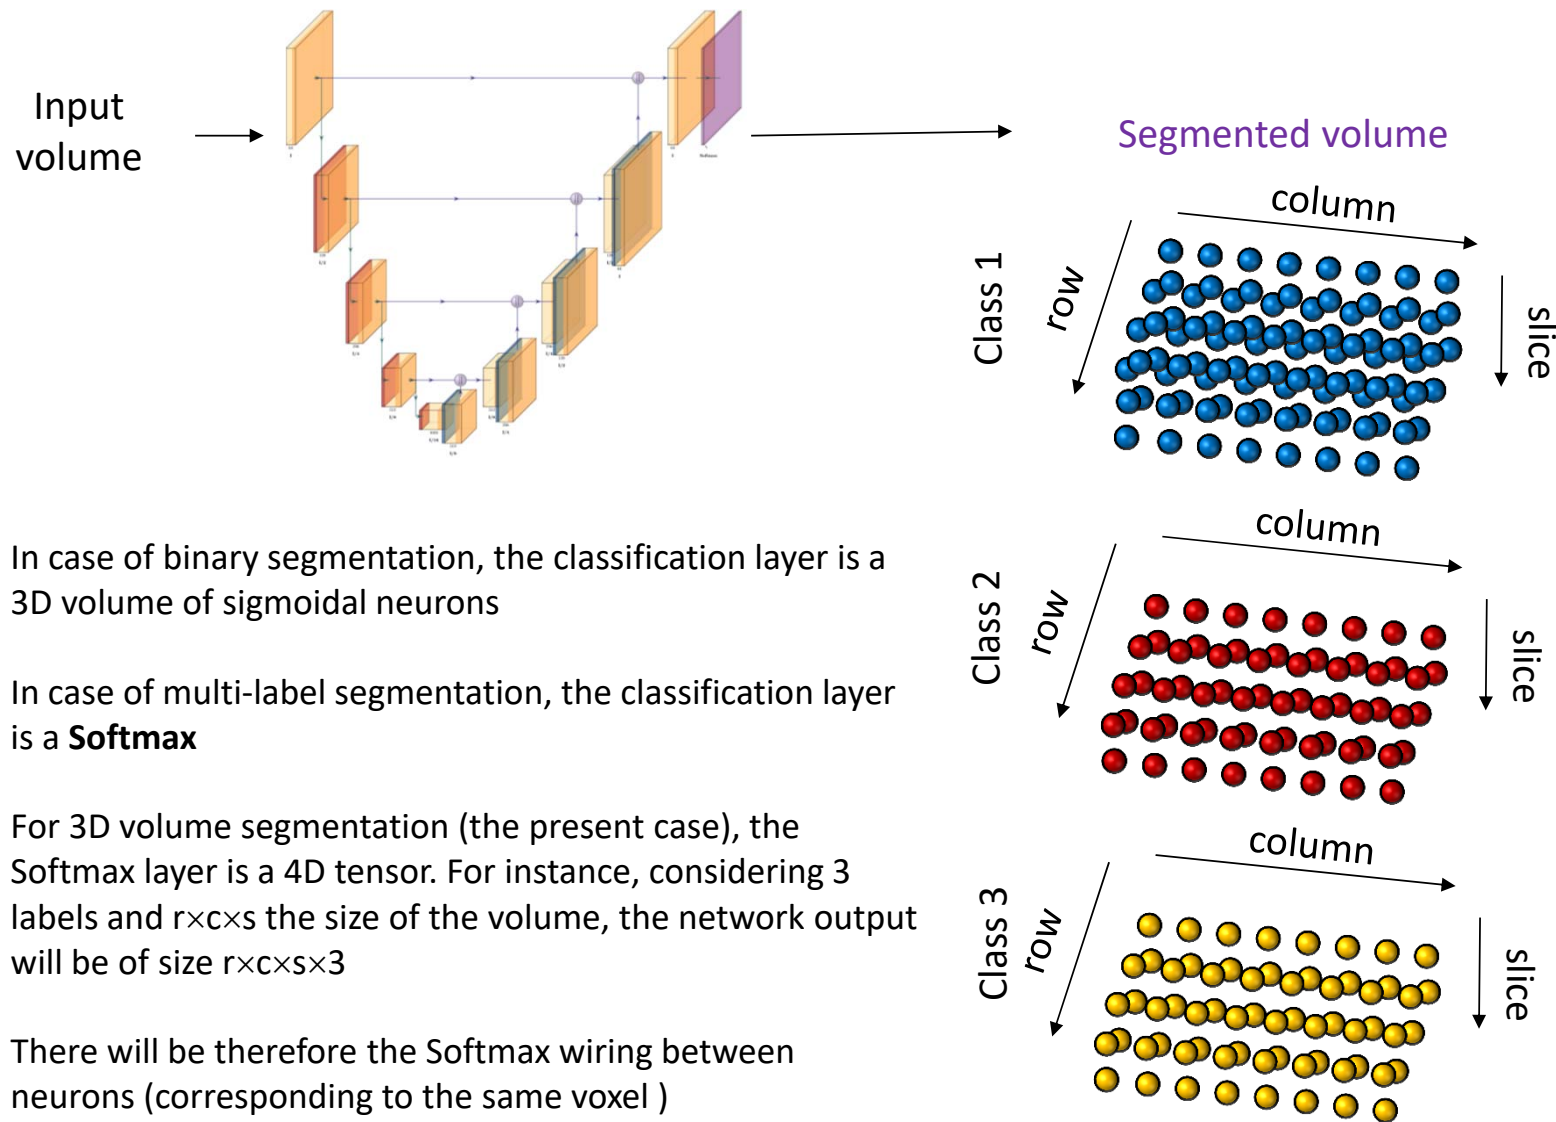


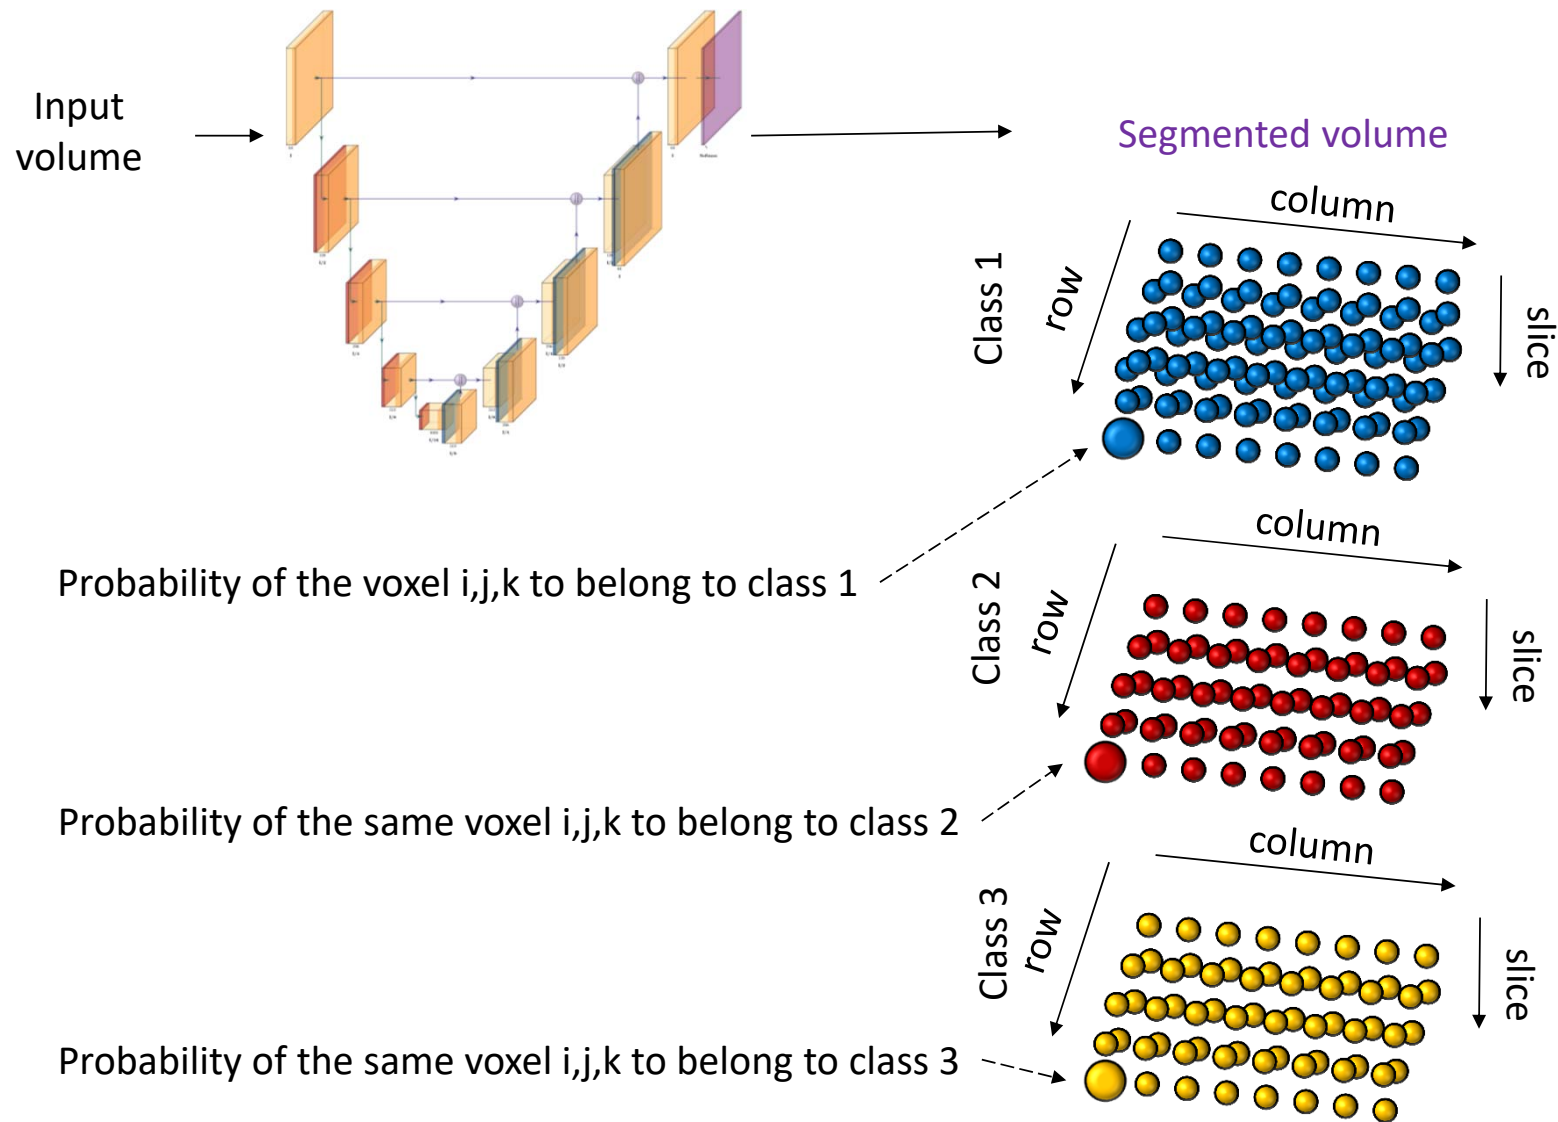
Output



U-NET for image segmentation

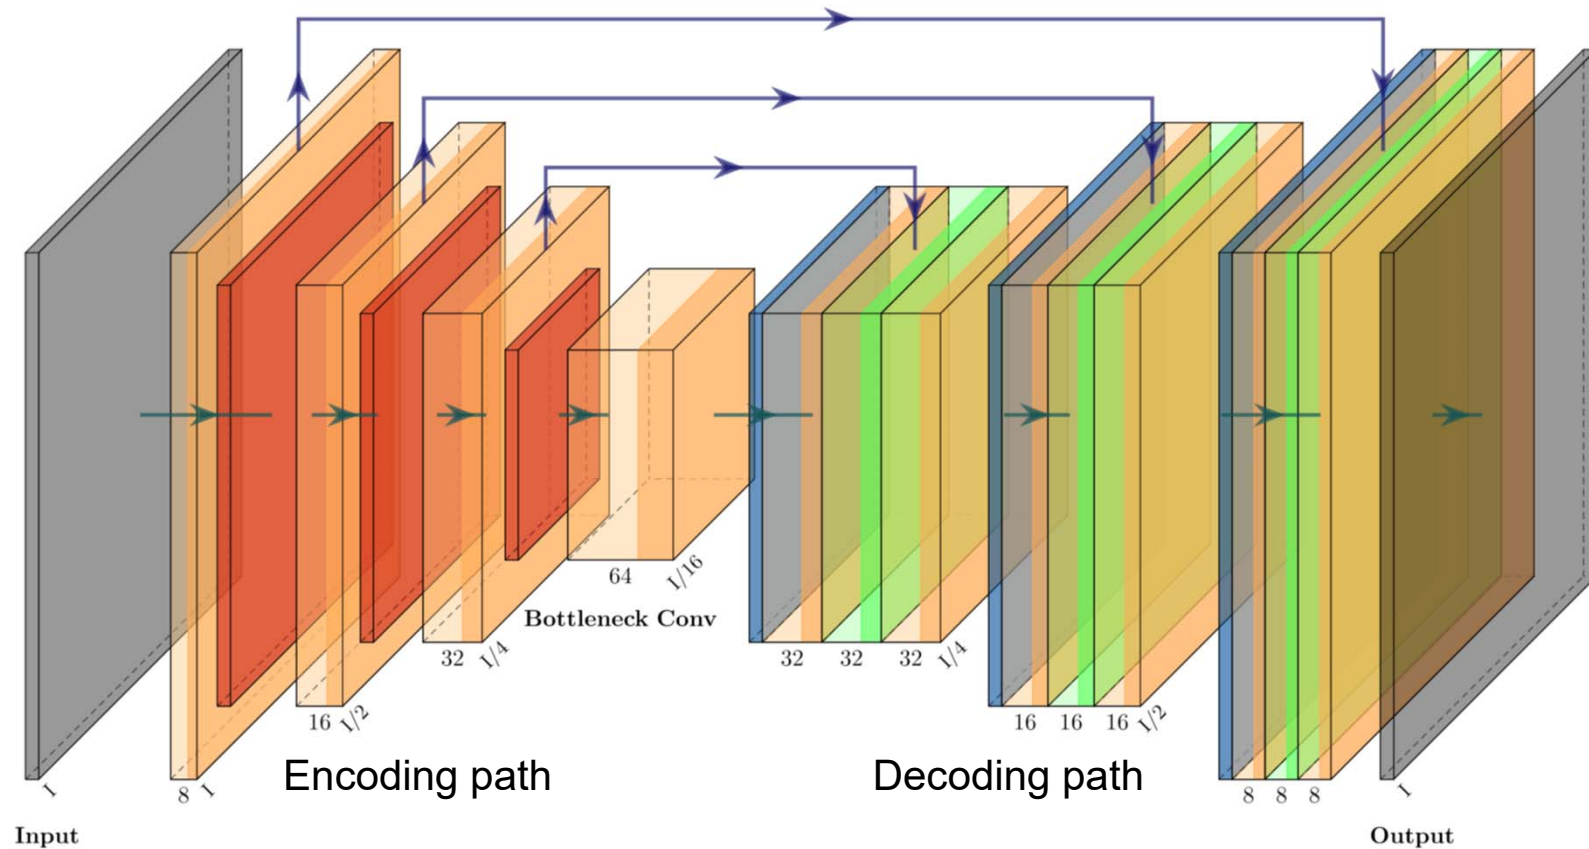






$$p1 + p2 + p3 = 1$$

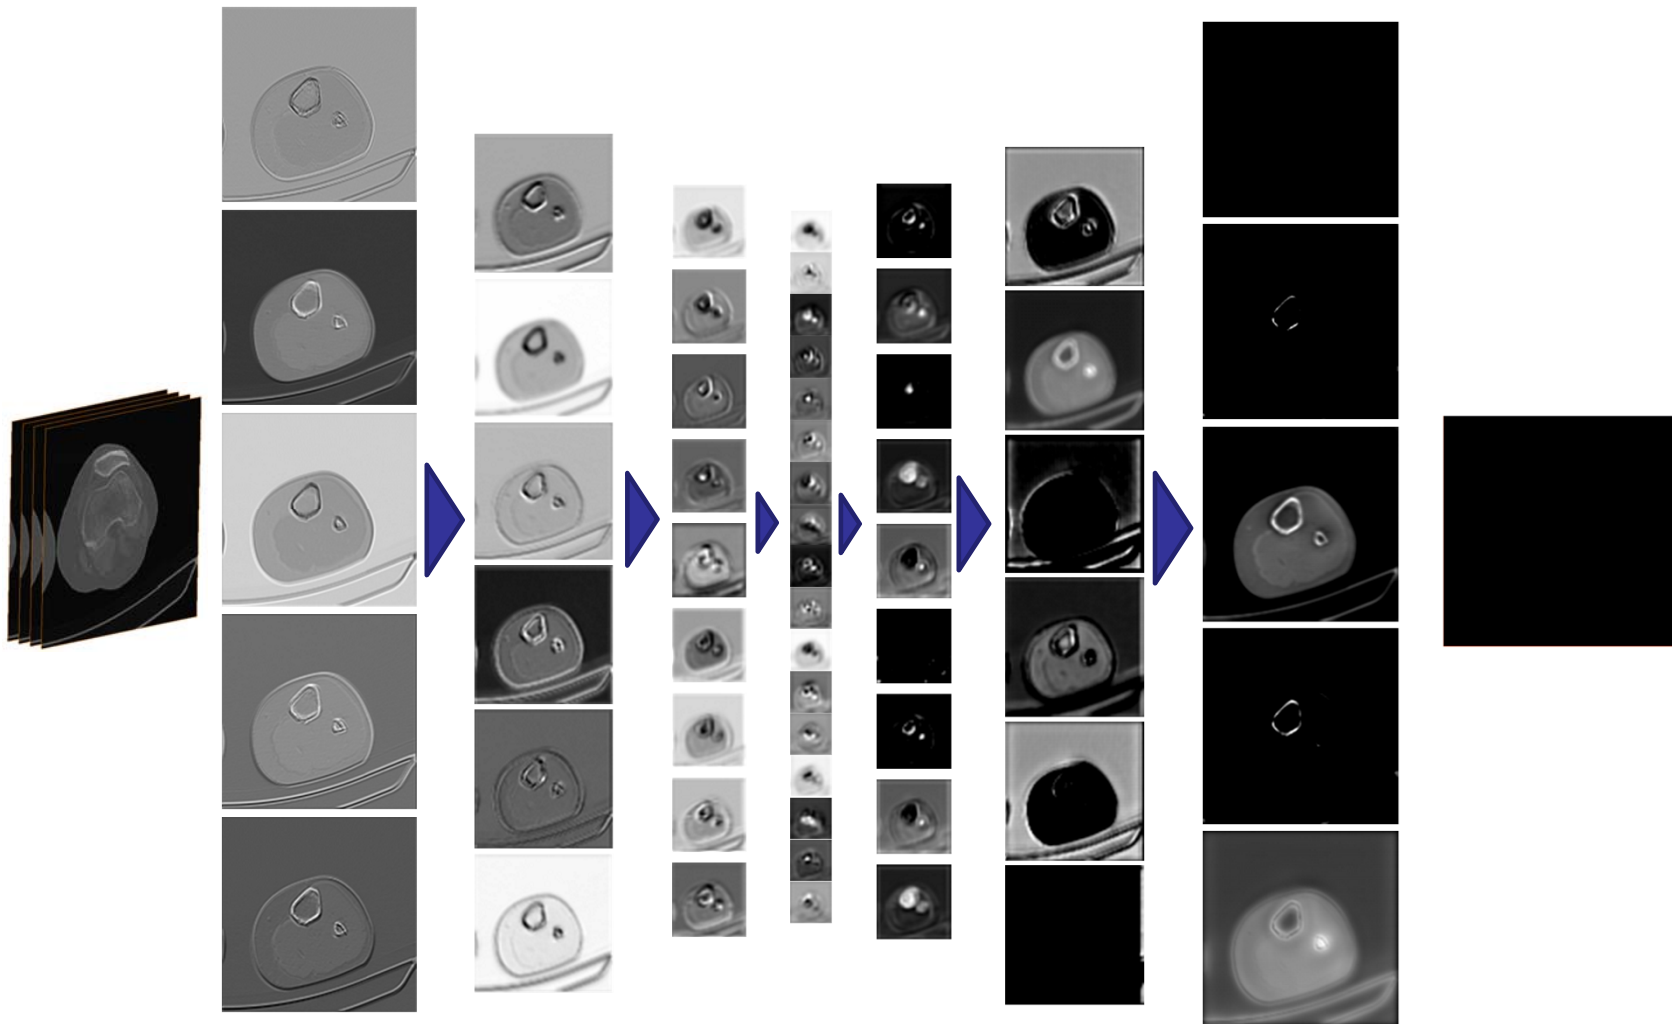
U-Net for tibia and femur segmentation in the knee



This network has 351435 parameters (filter size $3 \times 3 \times 3$)

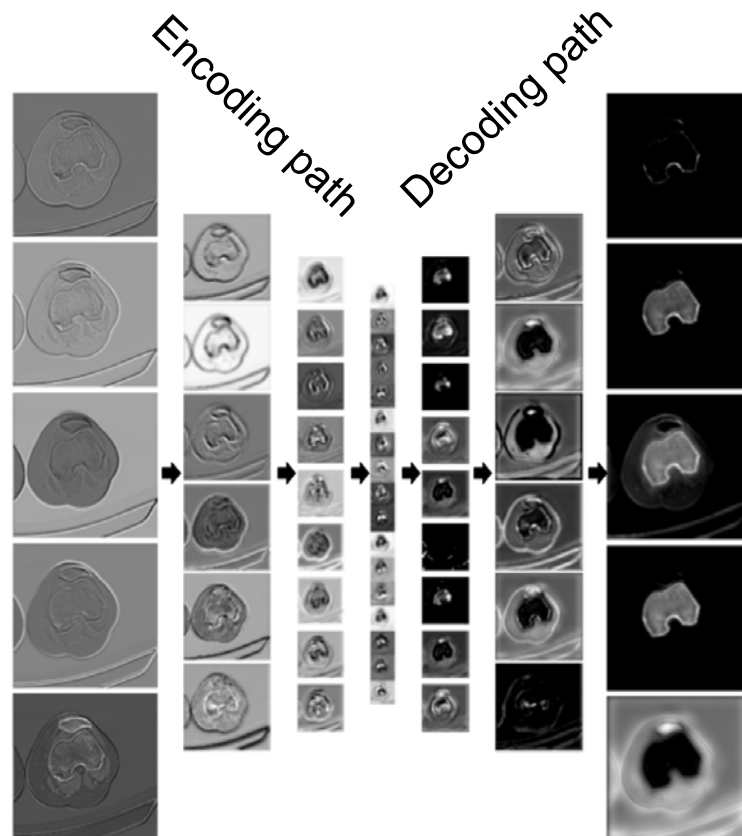
Assuming a volume of $128 \times 128 \times 128$ the network output will be $128 \times 128 \times 128 \times 3$

8-16-32-64-32-16-8



Relation with explainable AI

Explainable AI (XAI) investigates methods for analyzing or complementing AI models to make the internal logic and output of algorithms transparent and interpretable, making these processes humanly understandable and meaningful.



The encoding path is learning filters to extract geometrical features performing edge detection (easy part)

but... what about the decoding path??

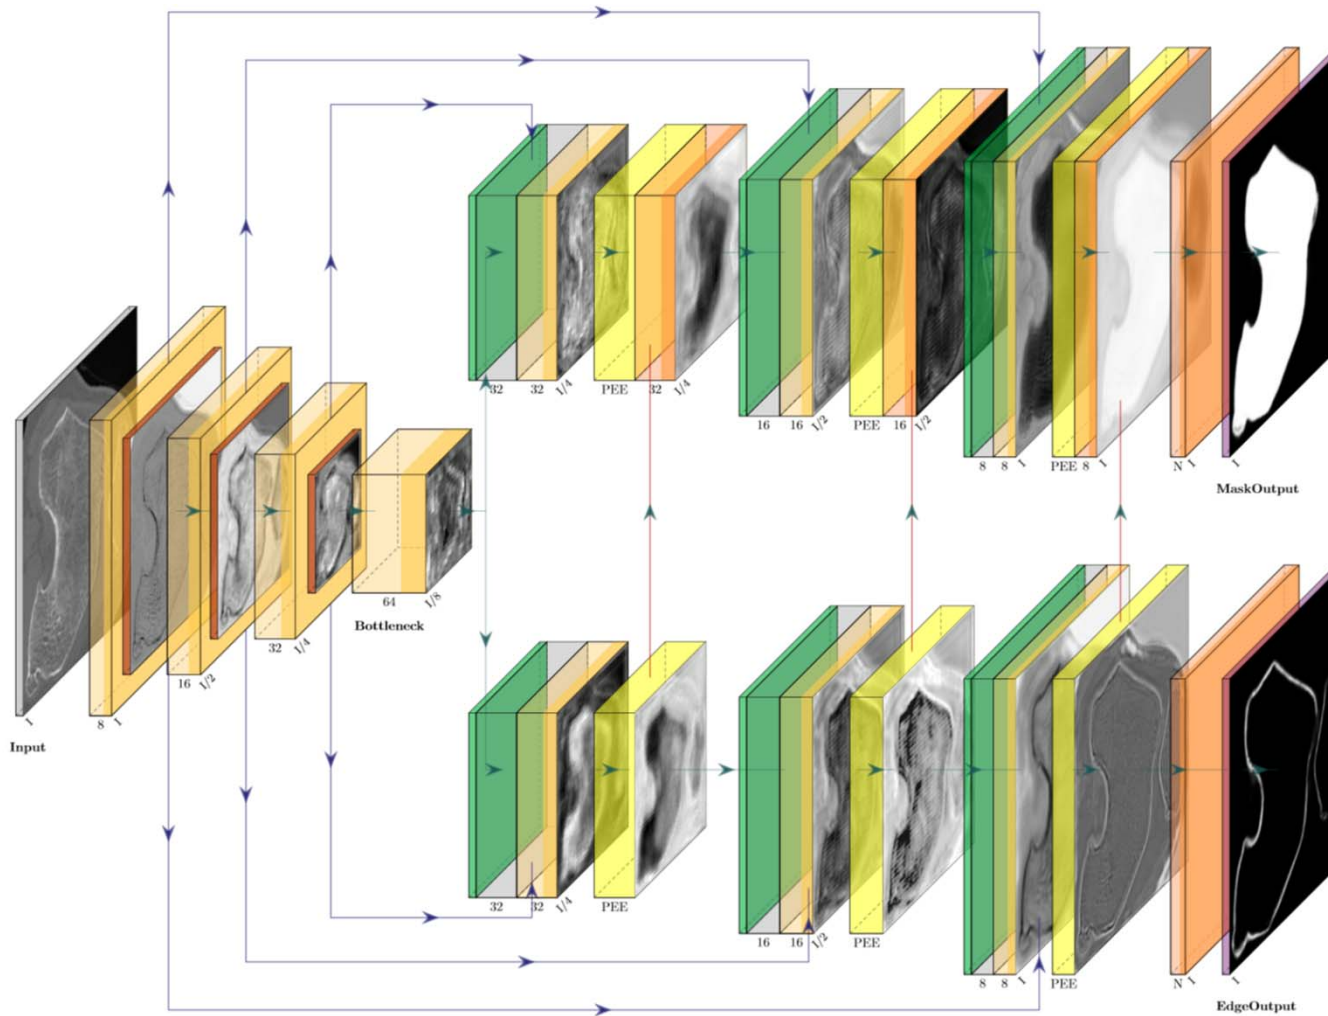
For sure it is learning semantics related to the task (femur and tibia segmentation)

Let us take into account that the 2D bone texture is very similar across all the 4 bones

So it is likely the decoding part is reconstructing the target shapes (2D/3D)

but.. what is the layer that has learnt to “remove” patella and fibula from segmentation?

Extending the U-NET with two interacting decoding paths




Upper decoding branch: masks
Lower decoding branch: contours

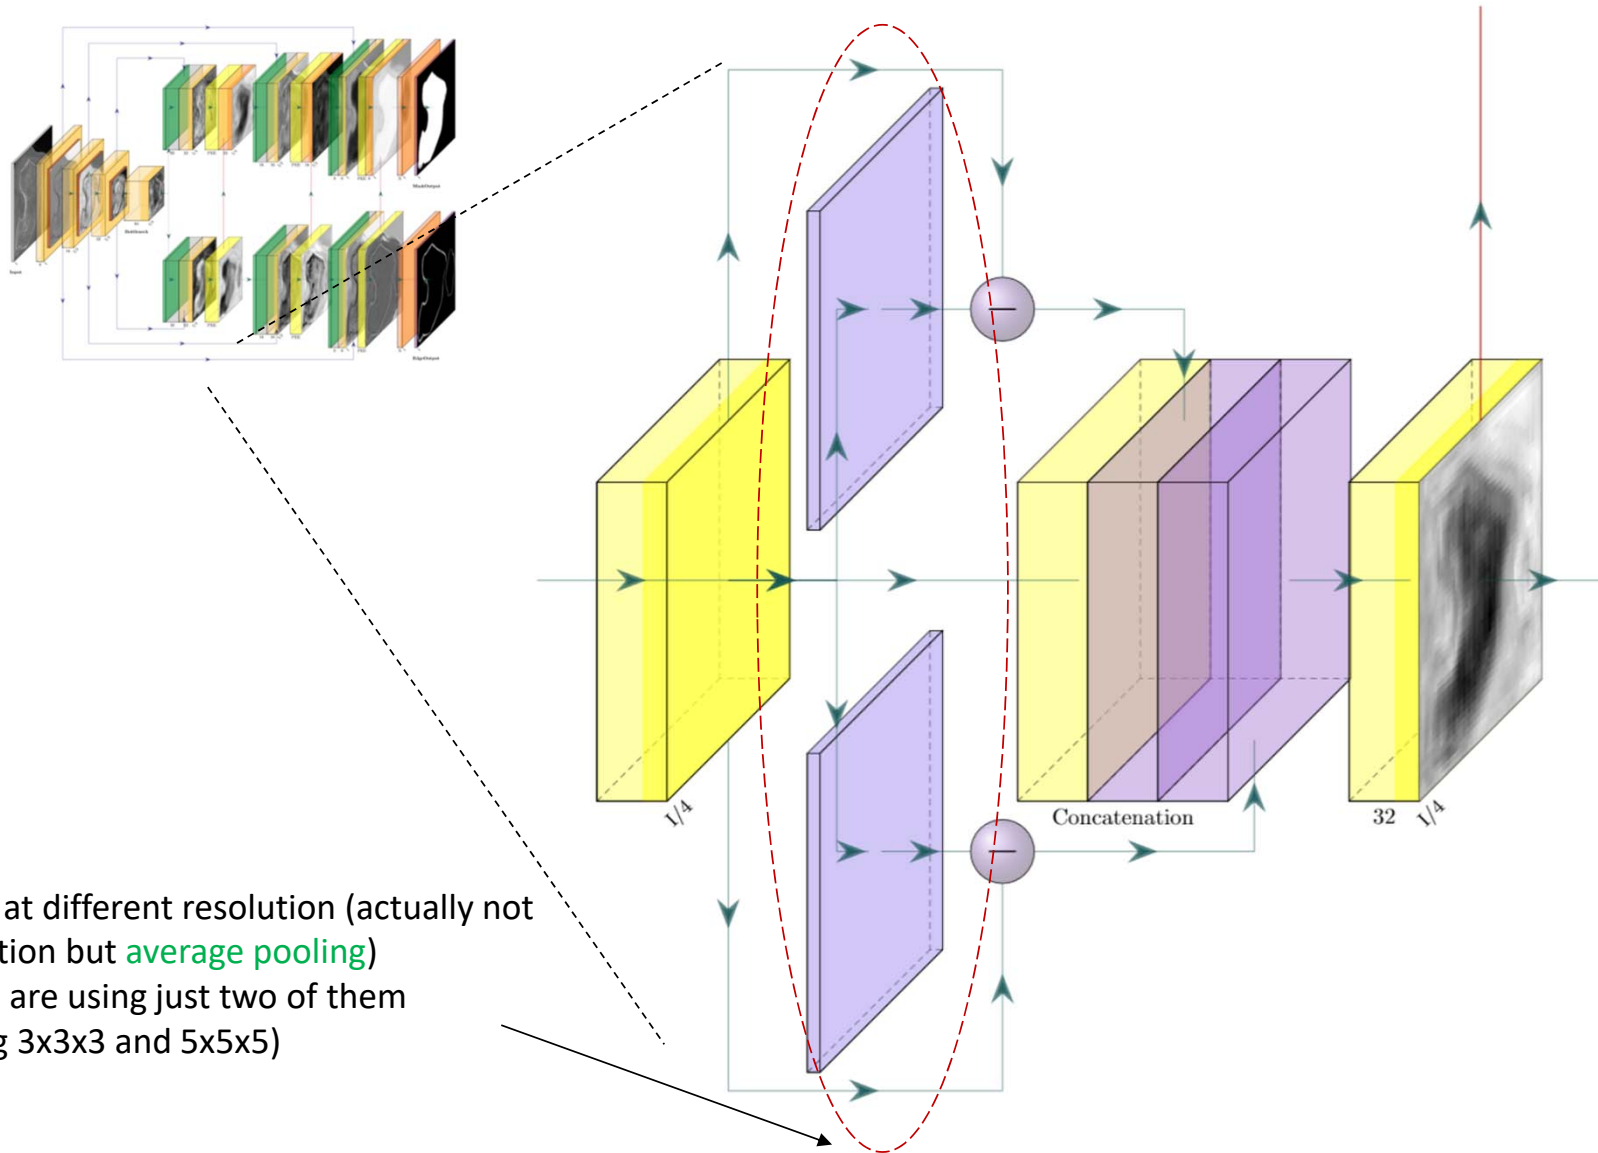
Skip connections (blue arrows)
Contour to mask concatenation (red arrows)

Pyramidal edge extraction (PEE) is implemented in each level of the edge detection branch.

Focus od pyramidal edge extraction module



Focus on pyramidal edge extraction (PEE) module




- ❑ “filters” at different resolution (actually not convolution but **average pooling**)
- ❑ here we are using just two of them
 - (e.g 3x3x3 and 5x5x5)

Focus od pyramidal edge extraction module

filters = 3x3x3

strides = 1



$x = \text{conv_layer}(\text{filters} / 2, \text{strides}, \text{padding}='same')(x)$

$x_1 = \text{average_layer}(\text{pool_size}=\text{pool_size}_1, \text{strides}=\text{strides}, \text{padding}='same')(x)$

$x_2 = \text{average_layer}(\text{pool_size}=\text{pool_size}_2, \text{strides}=\text{strides}, \text{padding}='same')(x)$


$x_{11} = \text{subtract}()([x, x_1])$

$x_{22} = \text{subtract}()([x, x_2])$

$x = \text{Concatenate}()([x, x_{11}, x_{22}])$


$x = \text{conv_layer}(\text{filters}, \text{strides}, \text{padding}='same')(x)$

Inception module: multi-resolution convolutional filters



Multiple kernel filter sizes on the same level

Inception module: multi-resolution convolutional filters and **dimensionality reduction**



DIFFERENCE?????